

Staged Compilation With Two-Level Type Theory

ANDRÁS KOVÁCS, Eötvös Loránd University, Hungary

The aim of staged compilation is to enable metaprogramming in a way such that we have guarantees about the well-formedness of code output, and we can also mix together object-level and meta-level code in a concise and convenient manner. In this work, we observe that two-level type theory (2LTT), a system originally devised for the purpose of synthetic homotopy theory, also serves as a system for staged compilation with dependent types. 2LTT has numerous good properties for this use case: it has a concise specification, well-developed algebraic and categorical model theory, and it supports a wide range of language features both at the object and the meta level. First, we give an overview of 2LTT's features and applications in staging. Then, we present a staging algorithm and provide a proof of correctness. Our algorithm is "staging-by-evaluation", analogously to the technique of normalization-by-evaluation, in that staging is given by the evaluation of 2LTT syntax in a semantic domain. The staging algorithm together with its correctness proof constitutes a proof of strong conservativity of 2LTT over the object theory. To our knowledge, this is the first system for staged compilation which supports full dependent types and unrestricted staging for types.

Additional Key Words and Phrases: type theory, two-level type theory, staged compilation

ACM Reference Format:

András Kovács. 2018. Staged Compilation With Two-Level Type Theory. *J. ACM* 37, 4, Article 111 (August 2018), 13 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 INTRODUCTION

The purpose of staged compilation is to write code-generating programs in a safe, ergonomic and expressive way. It is always possible to do ad-hoc code generation, by simply manipulating strings or syntax trees in a sufficiently expressive programming language. However, these approaches tend to suffer from verbosity, non-reusability and lack of safety. In staged compilation, there are certain *restrictions* on which metaprograms are expressible. Usually, staged systems enforce typing discipline, prohibit arbitrary manipulation of object-level scopes, and often they also prohibit accessing the internal structure of object expressions. On the other hand, we get *guarantees* about the well-scoping or well-typing of the code output, and we are also able to use concise syntax for embedding object-level code.

Two-level type theory, or 2LTT in short, was described by Annekov, Capriotti, Kraus and Sattler [1], building on ideas from Vladimir Voevodsky [3]. The motivation was to allow convenient metatheoretical reasoning about a certain mathematical language (homotopy type theory), and to enable concise and modular ways to extend the language with axioms.

It turns out that metamathematical convenience closely corresponds to metaprogramming convenience: 2LTT can be directly and effectively employed in staged compilation. Moreover, semantic ideas underlying 2LTT are also directly applicable to the theory of staging.

Author's address: András Kovács, kovacsandras@inf.elte.hu, Eötvös Loránd University, Hungary, Budapest.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Association for Computing Machinery.

0004-5411/2018/8-ART111 \$15.00

<https://doi.org/XXXXXXX.XXXXXXX>

1.1 Contributions

- In ?? we present an informal syntax of two-level type theory, a dependent type theory with staging features. We look at basic use-cases involving inlining control, partial evaluation and fusion optimizations. We also describe feature variations, enabling applications in monomorphization and memory layout control.
- In ??, following [1], we present a formal syntax of 2LTT and the object theory (the target theory of code generation). We recall the standard presheaf model of 2LTT, which lies over the syntactic category of the object theory. We show that the evaluation of 2LTT syntax in the presheaf model yields a staging algorithm.
- In ?? we show correctness of staging, consisting of
 - *Stability*: staging the output of staging has no action.
 - *Soundness*: the output of staging is convertible to the input.
 - *Completeness*: convertible programs produce convertible staging outputs.

Staging together with its correctness can be viewed as a *strong conservativity* theorem of 2LTT over the object theory. This means that the possible object-level constructions in 2LTT are in bijection with the constructions in the object theory, and staging witnesses that meta-level constructions can be always computed away. This improves on the weak notion of conservativity shown in [2] and [1].

- To our knowledge, this is the first description of a language which supports staging in the presence of full-blown dependent types, with universes and large elimination. Moreover, we allow unrestricted staging for types, so that types can be computed by metaprograms at compile time.

2 A TOUR OF TWO-LEVEL TYPE THEORY

In this section, we provide a short overview of 2LTT and its potential applications in staging. We work in the informal syntax of a dependently typed language which resembles Agda [?]. We focus on examples and informal explanations here; the formal details will be presented in Section [?].

Notation 1. We use the following notations throughout the paper. $(x : A) \rightarrow B$ denotes a dependent function type, where x may occur in B . We use $\lambda x. t$ for abstraction. A Σ -type is written as $(x : A) \times B$, with pairing as (t, u) , projections as fst and snd , and we may use pattern matching notation on pairs, e.g. as in $\lambda (x, y). t$. The unit type is \top with element tt . We will also use Agda-style notation for implicit arguments, where $t : \{x : A\} \rightarrow B$ implies that the first argument to t is inferred by default, and we can override this by writing a $t\{u\}$ application. We may also implicitly quantify over arguments (in the style of Idris and Haskell), for example when declaring $\text{id} : A \rightarrow A$ with the assumption that A is universally quantified.

2.1 Rules of 2LTT

Universes. We have universes $U_{i,j}$, where $i \in \{0, 1\}$, and $j \in \mathbb{N}$. The i index denotes stages, where 0 is the runtime (object-level) stage, and 1 is the compile time (meta-level) stage. The j index denotes universe sizes in the usual sense of type theory. We assume Russell-style universes, with $U_{i,j} : U_{i,j+1}$. However, for the sake of brevity we will usually omit the j indices in this section, since sizing is orthogonal to our use-cases and examples.

- U_0 can be viewed as the *universe of object-level or runtime types*. Each closed type $A : U_0$ can be staged to an actual type in the object language (the language of the staging output).
- U_1 can be viewed as the *universe of meta-level or static types*. If we have $A : U_1$, then A is guaranteed to be only present at compile time, and will be staged away. Elements $a : A$ are likewise computed away.

Type formers. U_0 and U_1 may be closed under arbitrary type formers, such as functions, Σ -types, identity types or inductive types in general. However, all constructors and eliminators in type formers must stay at the same stage. For example:

- Function domain and codomain types must be at the same stage.
- If we have $\text{Nat}_0 : U_0$ for the runtime type of natural numbers, we can only map from it to a type in U_0 by recursion or induction.

It is not required that we have the *same* type formers at both stages. We will discuss setups with different languages at different stages in Section ??.

Moving between stages. At this point, our system is rather limited, since there is no interaction between the stages. We enable such interaction through the following three operations.

- *Lifting:* for $A : U_0$, we have $\uparrow A : U_1$. From the staging point of view, $\uparrow A$ is the type of metaprograms which compute to runtime expressions of type A .
- *Quoting:* for $A : U_0$ and $t : A$, we have $\langle t \rangle : \uparrow A$. A quoted term $\langle t \rangle$ represents the metaprogram which immediately computes to t .
- *Splicing:* for $A : U_0$ and $t : \uparrow A$, we have $\sim t : A$. During staging, the metaprogram in the splice is executed, and the resulting expression is inserted into the output.

Notation 2. Splicing binds stronger than any operation, including function application. For instance, $\sim f x$ is parsed as $(\sim f) x$.

- Quoting and splicing are definitional inverses, i.e. we have $\sim \langle t \rangle = t$ and $\langle \sim t \rangle = t$ as definitional equalities.

Note that none of these three operations can be expressed as functions, since function types cannot cross between stages.

Informally, if we have a closed program $t : A$ with $A : U_0$, *staging* means computing all metaprograms and recursively replacing all splices in t and A with the resulting runtime expressions. The rules of 2LTT ensure that this is possible, and we always get a splice-free runtime program after staging.

Remark. Why do we use the index 0 for the runtime stage? The reason is that it is not difficult to generalize 2LTT to multi-level type theory, by allowing to lift types from U_i to U_{i+1} . In the semantics, this can be modeled by having a 2LTT whose object theory is once again a 2LTT, and doing this in an iterated fashion. But there must be necessarily a bottom-most object theory; hence our stage indexing scheme. For now though, we leave the multi-level generalization to future work.

Notation 3. We may disambiguate type formers at different stages by using 0 or 1 subscripts. For example, $\text{Nat}_1 : U_1$ is distinguished from $\text{Nat}_0 : U_0$, and likewise we may write $\text{zero}_0 : \text{Nat}_0$ and so on. For function and Σ types, the stage is usually easy to infer, so we do not annotate them. For example, the type $\text{Nat}_0 \rightarrow \text{Nat}_0$ must be at the runtime stage, since the domain and codomain types are at that stage, and we know that the function type former stays within a single stage. We may also omit stage annotations from λ and pairing.

2.2 Staged Programming in 2LTT

In 2LTT, we may have several different polymorphic identity functions. First, we consider the usual identity functions at each stage:

$$\begin{array}{ll} id_0 : (A : U_0) \rightarrow A \rightarrow A & id_1 : (A : U_1) \rightarrow A \rightarrow A \\ id_0 := \lambda A x. x & id_1 := \lambda A x. x \end{array}$$

An id_0 application will simply appear in staging output as it is. In contrast, id_1 can be used as a compile-time evaluated function, because the staging operations allow us to freely apply id_1 to runtime arguments. For example, $id_1 (\uparrow \text{Bool}_0) \langle \text{true}_0 \rangle$ has type $\uparrow \text{Bool}_0$, therefore $\sim(id_1 (\uparrow \text{Bool}_0) \langle \text{true}_0 \rangle)$ has type Bool_0 . We can stage this expression as follows:

$$\sim(id_1 (\uparrow \text{Bool}_0) \langle \text{true}_0 \rangle) = \sim \langle \text{true}_0 \rangle = \text{true}_0$$

There is another identity function, which computes at compile time, but which can be only used on runtime arguments:

$$\begin{aligned} id_{\uparrow} &: (A : \uparrow U_0) \rightarrow \uparrow \sim A \rightarrow \uparrow \sim A \\ id_{\uparrow} &:= \lambda A x. x \end{aligned}$$

Note that since $A : \uparrow U_0$, we have $\sim A : U_0$, hence $\uparrow \sim A : U_1$. Also, $\uparrow U_0 : U_1$, so all function domain and codomain types in the type of id_{\uparrow} are at the same stage. Now, we may write $\sim(id_{\uparrow} \langle \text{Bool}_0 \rangle \langle \text{true}_0 \rangle)$ for a term which is staged to true_0 . In this specific case id_{\uparrow} has no practical advantage over id_1 , but in some cases we really have to quantify over $\uparrow U_0$. This brings us to the next example.

Assume $\text{List}_0 : U_0 \rightarrow U_0$ with $\text{nil}_0 : (A : U_0) \rightarrow \text{List}_0 A$, $\text{cons}_0 : (A : U_0) \rightarrow A \rightarrow \text{List}_0 A$ and $\text{foldr}_0 : (A B : U_0) \rightarrow (A \rightarrow B \rightarrow B) \rightarrow B \rightarrow \text{List}_0 A \rightarrow B$. We define a map function which “inlines” its function argument:

$$\begin{aligned} \text{map} &: (A B : \uparrow U_0) \rightarrow (\uparrow \sim A \rightarrow \uparrow \sim B) \rightarrow \uparrow(\text{List}_0 \sim A) \rightarrow \uparrow(\text{List}_0 \sim B) \\ \text{map} &:= \lambda A B f \text{ as}. \langle \text{foldr}_0 \sim A \sim B (\lambda a \text{ bs}. \text{cons}_0 \sim B \sim (f \langle a \rangle) \text{ bs}) (\text{nil}_0 \sim B) \sim \text{as} \rangle \end{aligned}$$

This map function can be defined with quantification over $\uparrow U_0$ but not over U_1 , because List_0 expects type parameters in U_0 , and there is no generic way to convert from U_1 to U_0 . Now, assuming $- +_0 - : \text{Nat}_0 \rightarrow \text{Nat}_0 \rightarrow \text{Nat}_0$ and $ns : \text{List}_0 \text{Nat}_0$, we have the following staging behavior:

$$\begin{aligned} &\sim(\text{map} \langle \text{Nat}_0 \rangle \langle \text{Nat}_0 \rangle (\lambda n. \langle \sim n +_0 10 \rangle) \langle ns \rangle) \\ &= \sim \langle \text{foldr}_0 \sim \langle \text{Nat}_0 \rangle \sim \langle \text{Nat}_0 \rangle (\lambda a \text{ bs}. \text{cons}_0 \sim B \sim \langle \sim \langle a \rangle +_0 10 \rangle \text{ bs}) (\text{nil}_0 \sim \langle \text{Nat}_0 \rangle) \sim \langle ns \rangle \rangle \\ &= \text{foldr}_0 \text{Nat}_0 \text{Nat}_0 (\lambda a \text{ bs}. a +_0 10) (\text{nil}_0 \text{Nat}_0) ns \end{aligned}$$

By using meta-level functions and lifted types, we already have control over inlining. However, if we want to do more complicated meta-level computation, it is convenient to use recursion or induction on meta-level type formers. A classic example in staged compilation is the power function for natural numbers, which evaluates the exponent at compile time. We assume the iterator function $\text{iter}_1 : \{A : U_1\} \rightarrow \text{Nat}_1 \rightarrow (A \rightarrow A) \rightarrow A \rightarrow A$, and runtime multiplication as $- *_0 -$.

$$\begin{aligned} \text{exp} &: \text{Nat}_1 \rightarrow \uparrow \text{Nat}_0 \rightarrow \uparrow \text{Nat}_0 \\ \text{exp} &:= \lambda x y. \text{iter}_1 x (\lambda n. \langle \sim y *_0 \sim n \rangle) \langle 1 \rangle \end{aligned}$$

Now, $\sim(\text{exp} 3 \langle n \rangle)$ stages to $n *_0 n *_0 n *_0 1$ by the computation rules of iter_1 and the staging operations.

We can also stage *types*. Below, we use iteration to compute the type of vectors with static length, as a nested pair type.

$$\begin{aligned} \text{Vec} &: \text{Nat}_1 \rightarrow \uparrow U_0 \rightarrow \uparrow U_0 \\ \text{Vec} &:= \lambda n A. \text{iter}_1 n (\lambda B. \langle \sim A \times \sim B \rangle) \langle \top_0 \rangle \end{aligned}$$

With this definition, $\sim(\text{Vec} 3 \langle \text{Nat}_0 \rangle)$ stages to $\text{Nat}_0 \times (\text{Nat}_0 \times (\text{Nat}_0 \times \top_0))$. Now, we can use *induction* on Nat_1 to implement a map function. For readability, we use an Agda-style pattern

matching definition below (instead of the elimination principle).

$$\begin{aligned} \text{map} &: (n : \text{Nat}_1) \rightarrow (\uparrow \sim A \rightarrow \uparrow \sim B) \rightarrow \uparrow(\text{Vec } n A) \rightarrow \uparrow(\text{Vec } n B) \\ \text{map zero}_1 \quad f \text{ as} &:= \langle \text{tt}_0 \rangle \\ \text{map}(\text{suc}_1 n) f \text{ as} &:= \langle (\sim(f \langle \text{fst}_0 \sim \text{as} \rangle)), \text{map } n f \langle \text{snd}_0 \sim \text{as} \rangle \rangle \end{aligned}$$

This definition inlines the mapping function for each projected element of the vector. For instance, staging $\sim(\text{map } 2 (\lambda n. \langle \sim n +_0 10 \rangle) \langle \text{ns} \rangle)$ yields $(\text{fst}_0 \text{ ns} +_0 10, (\text{fst}_0(\text{snd}_0 \text{ ns}) +_0 10, \text{tt}_0))$. Sometimes, we do not want to duplicate the code of the mapping function. In such cases, we can use *let-insertion* [?], a standard technique in staged compilation. If we bind a runtime expression to a runtime variable, and only use that variable in subsequent staging, only the variable itself can be duplicated. One solution is to do an ad-hoc let-insertion:

$$\begin{aligned} \text{let}_0 f &:= \lambda n. n +_0 10 \text{ in } \sim(\text{map } 2 (\lambda n. \langle f \sim n \rangle) \langle \text{ns} \rangle) \\ &= \text{let}_0 f := \lambda n. n +_0 10 \text{ in } (f(\text{fst}_0 \text{ ns}), (f(\text{fst}_0(\text{snd}_0 \text{ ns})), \text{tt}_0)) \end{aligned}$$

Alternatively, we can define *map* so that it performs let-insertion, and we can switch between the two versions as needed.

More generally, we are free to use dependent types at the meta-level, so we can reproduce more complicated staging examples. Any well-typed interpreter can be rephrased as a *partial evaluator*, as long as we have sufficient type formers. For instance, we may write a partial evaluator for a simply typed lambda calculus. We sketch the implementation in the following. First, we inductively define contexts, types and terms:

$$\text{Ty} : \text{U}_1 \quad \text{Con} : \text{U}_1 \quad \text{Tm} : \text{Con} \rightarrow \text{Ty} \rightarrow \text{U}_1$$

Then we define the interpretation functions:

$$\begin{aligned} \text{EvalTy} &: \text{Ty} \rightarrow \uparrow \text{U}_0 \\ \text{EvalCon} &: \text{Con} \rightarrow \text{U}_1 \\ \text{EvalTm} &: \text{Tm } \Gamma A \rightarrow \text{EvalCon } \Gamma \rightarrow \uparrow \sim(\text{EvalTy } A) \end{aligned}$$

Types are necessarily computed to runtime types, e.g. an embedded representation of the natural number type is evaluated to $\langle \text{Nat}_0 \rangle$. Contexts are computed as follows:

$$\begin{aligned} \text{EvalCon empty} &:= \top_1 \\ \text{EvalCon (extend } \Gamma A) &:= \text{EvalCon } \Gamma \times (\uparrow \sim(\text{EvalTy } A)) \end{aligned}$$

This is a nice example for the usage of *partially static data*[?]: semantic contexts are *static* lists storing *runtime* expressions. This allows us to completely eliminate environment lookups in the staging output: an embedded lambda expression is staged to the corresponding lambda expression in the runtime language. This is similar to the partial evaluator presented in Idris 1 [?]. However, in contrast to 2LTT, Idris 1 does not provide a formal guarantee that partial evaluation does not get stuck.

2.3 Properties of Lifting & Binding Time Improvements

We describe more generally the action of \uparrow on type formers. First, \uparrow preserves negative type formers up to definitional isomorphism [1]:

$$\begin{aligned} \uparrow((x : A) \rightarrow B x) &\simeq ((x : \uparrow A) \rightarrow \uparrow(B \sim x)) \\ \uparrow((x : A) \times B) &\simeq ((x : \uparrow A) \times \uparrow(B \sim x)) \\ \uparrow \top_0 &\simeq \top_1 \end{aligned}$$

For function types, the preservation maps are the following:

$$\begin{aligned} pres_{\rightarrow} &: \uparrow((x : A) \rightarrow Bx) \rightarrow ((x : \uparrow A) \rightarrow \uparrow(B \sim x)) \\ pres_{\rightarrow} f &:= \lambda x. \langle f \sim x \rangle \\ pres_{\rightarrow}^{-1} f &:= \langle \lambda x. \sim(f \langle x \rangle) \rangle \end{aligned}$$

With this, we have that $pres_{\rightarrow}(pres_{\rightarrow}^{-1} f)$ is definitionally equal to f , and also the other way around. Preservation maps for Σ and \top work analogously.

By rewriting a 2LTT program left-to-right along preservation maps, we perform what is termed a *binding time improvement* in the partial evaluation literature [?]. Note that the output of $pres_{\rightarrow}$ uses a meta-level λ , while going the other way we introduce a runtime binder. Meta-level function and Σ types support more computation during staging, so in many cases it is beneficial to use the improved forms. In some cases though we may want to use unimproved forms, to limit the size of generated code. This is similar to what we have seen with let-insertion. For a minimal example, consider the following unimproved version of id_{\uparrow} :

$$\begin{aligned} id_{\uparrow} &: (A : \uparrow U_0) \rightarrow \uparrow(\sim A \rightarrow \sim A) \\ id_{\uparrow} &:= \lambda A. \langle \lambda x. x \rangle \end{aligned}$$

This can be used at the runtime stage as $\sim(id_{\uparrow} \langle Bool_0 \rangle) true_0$, which is staged to $(\lambda x. x) true_0$. This introduces a useless β -redex, so in this case the improved version is clearly preferable.

For inductive types in general we only get oplax preservation. For example, we have $Bool_1 \rightarrow \uparrow Bool_0$, defined as $\lambda b. \text{if } b \text{ then } \langle true_0 \rangle \text{ then } \langle false_0 \rangle$. In the staging literature, this is called “serialization” [?], or “lifting” in the context of Template Haskell [?]. In the other direction, we can only define constant functions from $\uparrow Bool_0$ to $Bool_1$.

The lack of elimination principles for $\uparrow A$ means that we cannot inspect the internal structure of expressions. This is called “generativity” in staging [?]. We will briefly discuss non-generative staging in Section ??.

In particular, we have no serialization map from $Nat_1 \rightarrow Nat_1$ to $\uparrow(Nat_0 \rightarrow Nat_0)$. However, when $A : U_1$ is *finite*, and $B : U_1$ can be serialized, then $A \rightarrow B$ can be serialized, because it is equivalent to a finite product. For instance, $Bool_1 \rightarrow Nat_1 \simeq Nat_1 \times Nat_1$. In 2LTT, A is called *cofibrant* [?]: this means that for each B , $A \rightarrow \uparrow B$ is equivalent to $\uparrow C$ for some C . This formalizes the so-called “trick” in partial evaluation, which improves binding times by η -expanding functions out of finite sums [?].

2.3.1 Fusion. Fusion optimizations can be viewed as binding time improvement techniques for general inductive types. The basic idea is that by lambda-encoding an inductive type, it is brought to a form which can be binding-time improved. For instance, consider foldr-build fusion for lists, which is employed in GHC Haskell [?]. Starting from $\uparrow(List_0 A)$, we use Böhm-Berarducci encoding [?] under the lifting to get

$$\uparrow((L : U_0) \rightarrow (A \rightarrow L \rightarrow L) \rightarrow L \rightarrow L)$$

which is isomorphic to

$$(L : \uparrow U_0) \rightarrow (\uparrow A \rightarrow \uparrow \sim L \rightarrow \uparrow \sim L) \rightarrow \uparrow \sim L \rightarrow \uparrow \sim L.$$

Alternatively, for *stream fusion*, we embed $List A$ into the coinductive colists (i.e. the possibly infinite lists), and use a terminal lambda-encoding. The embedding into the “larger” structure enables some staged optimizations which are otherwise not possible, such as fusion for the zip function [?]. However, the price we pay is that converting back to lists from colists is not necessarily total.

We do not detail the implementation of fusion in 2LTT here. In short, 2LTT is a natural setting for a wide range of fusion setups. A major advantage of fusion in 2LTT is the formal guarantee of staging, in contrast to implementations where compile-time computation relies on ad-hoc user annotations and general-purpose optimization passes. For instance, fusion in GHC relies on rewrite rules and inlining annotations which have to be carefully tuned and ordered, and it is quite possible to get pessimized code via failed fusion.

2.3.2 Inferring staging operations. During bidirectional elaboration [?], we can use the preservation isomorphisms and the quote-splice isomorphism as a coercive subtyping system. When elaboration needs to compare an inferred and an expected type, it may insert transports along isomorphisms. We implemented this feature in our prototype. It additionally supports Agda-style implicit arguments and pattern unification, so it can elaborate the following definition:

$$\begin{aligned} \text{map} &: \{AB : \uparrow U_0\} \rightarrow (\uparrow A \rightarrow \uparrow B) \rightarrow \uparrow(\text{List}_0 A) \rightarrow \uparrow(\text{List}_0 B) \\ \text{map} &:= \lambda f \text{ as. foldr}_0(\lambda a \text{ bs. cons}_0(f a) \text{ bs}) \text{ nil}_0 \text{ as} \end{aligned}$$

We may go a bit further, and also add the coercive subtyping rule $U_0 \leq U_1$, witnessed by \uparrow . Then, the type of *map* can be written as $\{AB : \uparrow U_0\} \rightarrow (A \rightarrow B) \rightarrow \text{List}_0 A \rightarrow \text{List}_0 B$. However, here the elaborator has to make a choice, whether to elaborate to improved or unimproved types. In this case, the fully unimproved type would be

$$\{AB : \uparrow U_0\} \rightarrow \uparrow((\sim A \rightarrow \sim B) \rightarrow \text{List}_0 \sim A \rightarrow \text{List}_0 \sim B).$$

It seems to the author of this paper that improved types are a sensible default, and we can insert explicit lifting when we want to opt for unimproved types. This is also available in our prototype.

2.4 Variations of Object-Level Languages

In the following, we consider variations on object-level languages, with a focus on applications in downstream compilation after staging. Adding restrictions or more distinctions to the object language can make it easier to optimize and compile it.

2.4.1 Monomorphization. In this case, the object language is simply typed, so every type is known statically. This makes it easy to assign different memory layouts to different types, and generate code accordingly for each type. Moving to 2LTT, we still want to abstract over runtime types at compile time, so we use the following setup.

- We have a *judgment*, written as $A \text{ type}_0$, for well-formed runtime types. Runtime types may be closed under simple type formers.
- We have a type $Ty_0 : U_1$ in lieu of the previous $\uparrow U_0$.
- For each $A \text{ type}_0$, we have $\uparrow A : Ty_0$.
- We have quoting and splicing for types and terms. For types, we send $A \text{ type}_0$ to $\langle A \rangle : Ty_0$. For terms, we send $t : A$ to $\langle t \rangle : \uparrow A$.

Despite the restriction to simple types at runtime, we can still write arbitrary higher-rank polymorphic functions in this setup, such as a function with type $((A : Ty_0) \rightarrow \uparrow \sim A \rightarrow \uparrow \sim A) \rightarrow \uparrow \text{Bool}_0$. This function can be only applied to statically known arguments in runtime code, so the polymorphism can be staged away. The main restriction that programmers have to keep in mind is that polymorphic functions cannot be stored inside runtime data types.

This setup is fairly similar to the well-known monomorphization models in the C++ and Rust programming languages. The evident difference is that in 2LTT we can use a full type theory at compile time, as opposed to languages which are less expressive and less principled. Additionally, we will see in Section ?? that this setup is compatible with an *induction principle* for Ty_0 , so that we can analyze the structure of runtime types.

2.4.2 Memory representation polymorphism. This refines monomorphization, in that types are not directly identified with memory representations, but instead representations are internalized in 2LTT as a meta-level type, and runtime types are indexed over representations.

- We have $\text{Rep} : \mathcal{U}_1$ as the type of memory representations. We have considerable freedom in the specification of Rep . A simple setup may distinguish references from unboxed products, i.e. we have $\text{Ref} : \text{Rep}$ and $\text{Prod} : \text{Rep} \rightarrow \text{Rep} \rightarrow \text{Rep}$, and additionally we may assume any desired primitive machine representation as a Rep .
- We have Russell-style $\mathcal{U}_{0,j} : \text{Rep} \rightarrow \mathcal{U}_{0,j+1} r$, where r is some chosen runtime representation for types; often this would mark types as erased at runtime. We leave the meta-level $\mathcal{U}_{1,j}$ hierarchy unchanged.
- We may introduce unboxed Σ types and primitive machine types in the runtime language. For $A : \mathcal{U}_0 r$ and $B : A \rightarrow \mathcal{U}_0 r'$, we may have $(x : A) \times Bx : \mathcal{U}_0 (\text{Prod } rr')$. Thus, we have type dependency, but we do not have dependency in memory representations.

Since Rep is meta-level, there is no way to abstract over it at runtime, and during staging all Rep indices are computed to concrete canonical representations. This is a way to reconcile dependent types with control over memory layouts. The unboxed flavor of Σ ends up with a statically known flat memory representation, computed from the representations of the fields.

In principle, it should be possible to have *dependent memory layouts*; this would be a more advanced variation on type-passing polymorphism [?], where layouts of objects may depend on runtime data. For example, length-prefixed flat arrays would not be a primitive notion, but simply defined as $(n : \text{Nat}) \times \text{Vec } n A$. However, it seems to be highly challenging to implement runtime systems and precise garbage collection generically over dependently typed memory layouts. See sixteen [?] for an experimental implementation with conservative garbage collection.

3 FORMAL 2LTT

In this section we switch to a formal description of 2LTT, which we will use in the subsequent sections to define the staging algorithm and prove its correctness. First we describe the metatheory that we work in.

3.1 Metatheory

TODO

3.2 Models and Syntax of 2LTT

We use an algebraic specification for models, and specify the syntax as the initial model. We only handle well-formed syntactic objects and only define operations which respect syntactic definitional equality; in fact, we identify definitional equality with metatheoretic equality. First, we define the structural scaffolding of 2LTT without type formers.

Definition 3.1. A model of **basic 2LTT** consists of the following.

- A category \mathbb{C} with a terminal object. We denote the set of objects as $\text{Con}_{\mathbb{C}} : \text{Set}$ and use capital Greek letters starting from Γ to refer to objects. The set of morphisms is $\text{Sub}_{\mathbb{C}} : \text{Con}_{\mathbb{C}} \rightarrow \text{Con}_{\mathbb{C}} \rightarrow \text{Set}$, and we use σ, δ and so on to refer to morphisms. The terminal object is written as \bullet with unique morphism $\epsilon : \text{Sub}_{\mathbb{C}} \Gamma \bullet$. We omit the \mathbb{C} subscript if it is clear from context.
- For each $i \in \{0, 1\}$ and $j \in \mathbb{N}$, we have $\text{Ty}_{i,j} : \text{Con} \rightarrow \text{Set}$ and $\text{Tm}_{i,j} : (\Gamma : \text{Con}) \rightarrow \text{Ty}_{i,j} \Gamma \rightarrow \text{Set}$, where Ty is a presheaf over \mathbb{C} and $\text{Tm}_{i,j}$ is a presheaf over the category of elements of $\text{Ty}_{i,j}$. This means that both types (Ty) and terms (Tm) can be substituted, and substitution has functorial action. We use A, B, C to refer to types and t, u, v to refer to terms, and use $A[\sigma]$ and $t[\sigma]$ for substituting types and terms. Additionally, for each $\Gamma : \text{Con}$ and

$A : \text{Ty}_{i,j} \Gamma$, we have the extended object $\Gamma \triangleright A : \text{Con}$ such that there is a natural isomorphism $\text{Sub } \Gamma (\Delta \triangleright A) \simeq ((\sigma : \text{Sub } \Gamma \Delta) \times \text{Tm}_{i,j} \Gamma (A[\sigma]))$.

- For each j we have a *lifting structure*, consisting of a natural transformation $\uparrow : \text{Ty}_{0,j} \Gamma \rightarrow \text{Ty}_{1,j} \Gamma$, and an invertible natural transformation $\langle - \rangle : \text{Tm}_{0,j} \Gamma A \rightarrow \text{Tm}_{1,j} \Gamma (\uparrow A)$, with inverse $\sim -$.

In short, for each i and j we have a family structure in the sense of categories-with-families [?], such that there is a family morphism from each $(\text{Ty}_{0,j}, \text{Tm}_{0,j})$ to $(\text{Ty}_{1,j}, \text{Tm}_{1,j})$ in the sense of [?]. The following notions are derivable in any model:

- By moving left-to-right along $\text{Sub } \Gamma (\Delta \triangleright A) \simeq ((\sigma : \text{Sub } \Gamma \Delta) \times \text{Tm}_{i,j} \Gamma (A[\sigma]))$, and starting from the identity morphism $\text{id} : \text{Sub } (\Gamma \triangleright A) (\Gamma \triangleright A)$, we recover the *weakening substitution* $p : \text{Sub } (\Gamma \triangleright A) \Gamma$ and the *zero variable* $q : \text{Tm}_{i,j} (\Gamma \triangleright A) (A[p])$.
- By weakening q , we recover a notion of variables as De Bruijn indices. In general, the n -th De Bruijn index is defined as $q[p^n]$, where p^n denotes n -fold composition.
- By moving right-to-left along $\text{Sub } \Gamma (\Delta \triangleright A) \simeq ((\sigma : \text{Sub } \Gamma \Delta) \times \text{Tm}_{i,j} \Gamma (A[\sigma]))$, we recover the operation which extends a morphism with a term. In the syntax (initial model), this justifies the view of Sub as a list of terms, i.e. a parallel substitution. We denote the extension operation as (σ, t) for $\sigma : \text{Sub } \Gamma \Delta$.

Notation 4. De Bruijn indices are rather hard to read, so we will often use nameful notation for binders and substitutions. For example, we may write $\Gamma \triangleright (x : A) \triangleright (y : B)$ for a context, and subsequently write $B[y \mapsto t]$ for substituting the x variable for some term $t : \text{Tm}_{i,j} \Gamma A$. Using nameless notation, we would instead have $B : \text{Ty}_{i,j} (\Gamma \triangleright A)$ and $B[\text{id}, t]$; here we recover single substitution by extending the identity substitution $\text{id} : \text{Sub } \Gamma \Gamma$ with t .

We may also use implicit weakening: if a type or term is in a Γ context, we may use it in an extended $\Gamma \triangleright A$ context without marking the weakening substitution.

Definition 3.2. A **model of 2LTT** is a model of basic 2LTT which supports certain type formers. For the sake of brevity, we only present our results for a small collection of type formers. However, we will argue that our results easily extend to any general notion of inductive type formers. We specify type formers in the following. We assume that all type formers are natural, i.e. stable under substitution.

- *Universes.* For each i and j , we have a Coquand-style universe [?] in $\text{Ty}_{i,j}$. This consists of $\text{U}_{i,j} : \text{Ty}_{i,j+1} \Gamma$, together with $\text{El} : \text{Tm}_{i,j+1} \Gamma \text{U}_{i,j} \rightarrow \text{Ty}_{i,j} \Gamma$ and Code , where Code and El are inverses.
- *Σ -types.* We have $\Sigma (x : A) B : \text{Ty}_{i,j} \Gamma$ for $A : \text{Ty}_{i,j} \Gamma$ and $B : \text{Ty}_{i,j} (\Gamma \triangleright (x : A))$, together with a natural isomorphism of pairing and projections:

$$\text{Tm}_{i,j} \Gamma (\Sigma A B) \simeq ((t : \text{Tm}_{i,j} \Gamma A) \times \text{Tm}_{i,j} \Gamma (B[x \mapsto t]))$$

We write (t, u) for pairing and fst and snd for projections.

- *Function types.* We have $\Pi (x : A) B : \text{Ty}_{i,j} \Gamma$ for $A : \text{Ty}_{i,j} \Gamma$ and $B : \text{Ty}_{i,j} (\Gamma \triangleright (x : A))$, together with $\text{app} : \text{Tm}_{i,j} \Gamma (\Pi A B) \rightarrow \text{Tm}_{i,j} \Gamma (\Gamma \triangleright (x : A)) B$ and its inverse lam .

- *Natural numbers.* We have $\text{Nat}_{i,j} : \text{Ty}_{i,j} \Gamma$, $\text{zero}_{i,j} : \text{Tm}_{i,j} \Gamma \text{Nat}_{i,j}$, and $\text{suc}_{i,j} : \text{Tm}_{i,j} \Gamma \text{Nat}_{i,j} \rightarrow \text{Tm}_{i,j} \Gamma \text{Nat}_{i,j}$. The eliminator is the following.

$$\begin{aligned} \text{NatElim} : & (P : \text{Ty}_{i,k} (\Gamma \triangleright (n : \text{Nat}_{i,j})) \\ & (z : \text{Tm}_{i,k} \Gamma (P[n \mapsto \text{zero}_{i,j}]))) \\ & (s : \text{Tm}_{i,k} (\Gamma \triangleright (n : \text{Nat}_{i,j}) \triangleright (pn : P[n \mapsto n])) (P[n \mapsto \text{suc}_{i,j} n])) \\ & (t : \text{Tm}_{i,j} \Gamma \text{Nat}_{i,j}) \\ \rightarrow & \text{Tm}_{i,j} \Gamma (P[n \mapsto t])) \end{aligned}$$

We also have the β -rules:

$$\begin{aligned} \text{NatElim } P \ z \ s \ \text{zero}_{i,j} &= z \\ \text{NatElim } P \ z \ s \ (\text{suc}_{i,j} \ t) &= s[n \mapsto t, pn \mapsto \text{NatElim } P \ z \ s \ t] \end{aligned}$$

Note that we can eliminate from a j level to any k level.

Definition 3.3. The **syntax of 2LTT** is defined to be the initial model of 2LTT. The syntax can be assumed to exist, since the notion of model is algebraic; the specification can be expressed using a suitable notion of algebraic signatures, such as one for essentially algebraic theories [?] or as a finitary quotient inductive-inductive signature [?]. The initiality of syntax directly yields a notion of *recursion on syntax*, and it can be shown that initiality also implies an *induction principle* [?].

The examples in Section ?? can be faithfully represented in the formal syntax. The only notable difference is that the formal syntax uses Coquand-style universes instead of Russell-style ones. We use the former because it is much easier to model in the semantics in Section ?. We note though that it is straightforward to elaborate from a Russell-style surface syntax to the formal version, by inserting `El` and `Code` as required.

3.2.1 Comparison to Annekov et al. Comparing our models to the primary reference on 2LTT [1], the main difference is the handling of “sizing” levels. In *ibid.* there is a cumulative lifting from $\text{Ty}_{i,j}$ to $\text{Ty}_{i,j+1}$, which we do not assume. Instead, we allow elimination from $\text{Nat}_{i,j}$ into any k level. This means that we can manually define lifting maps from $\text{Nat}_{i,j}$ to $\text{Nat}_{i,j+1}$ by elimination. This is more similar to e.g. Agda, where we do not have cumulativity, but we can define explicit lifting from $\text{Nat}_j : \text{Set}_j$ to $\text{Nat}_k : \text{Set}_k$.

In [1], “two-level type theory” specifically refers to the setup where the object-level is a homotopy type theory and the meta level is an extensional type theory. In contrast, we allow a wider range of setups under the 2LTT umbrella. Annekov et al. also considers a range of additional strengthenings and extension of 2LTT [1, Section 2.4], most of which are of use in synthetic homotopy theory. We do not assume any of these, and stick to the most basic formulation of 2LTT.

3.3 Models and Syntax of the Object Theory

We also need to specify the object theory, which serves as the output language of staging. In general, the object theory corresponding to a particular flavor of 2LTT is simply the type theory that supports only the object-level $\text{Ty}_{0,j}$ hierarchy with its type formers.

Definition 3.4. A **model of the object theory** is a category-with-families, with types and terms indexed over $j \in \mathbb{N}$, supporting Coquand-style universes U_j , type formers Π , Σ and Nat_j , with elimination from Nat_j to any k level.

Definition 3.5. Like before, the **syntax of the object theory** is the initial model.

Notation 5. From now on, by default we use Con , Sub , Ty and Tm to refer to sets in the syntax of 2LTT. We use Con_0 , Sub_0 , Ty_0 and Tm_0 to refer to its underlying sets.

Definition 3.6. We have an **embedding** of object syntax into 2LTT syntax, consisting of the following functions:

$$\begin{aligned} \ulcorner - \urcorner &: \text{Con}_0 \rightarrow \text{Con} \\ \ulcorner - \urcorner &: \text{Sub}_0 \Gamma \Delta \rightarrow \text{Sub} \ulcorner \Gamma \urcorner \ulcorner \Delta \urcorner \\ \ulcorner - \urcorner &: \text{Ty}_{0,j} \Gamma \rightarrow \text{Ty}_{0,j} \ulcorner \Gamma \urcorner \\ \ulcorner - \urcorner &: \text{Tm}_{0,j} \Gamma A \rightarrow \text{Tm}_{0,j} \ulcorner \Gamma \urcorner \ulcorner A \urcorner \end{aligned}$$

Embedding maps all type and term formers to the corresponding ones in 2LTT, and strictly preserves all structure.

4 THE STAGING ALGORITHM

In this section we specify what we mean by staging, then define a staging algorithm.

Definition 4.1. A **staging algorithm** consists of two functions:

$$\begin{aligned} \text{Stage} &: \text{Ty}_{0,j} \ulcorner \Gamma \urcorner \rightarrow \text{Ty}_{0,j} \Gamma \\ \text{Stage} &: \text{Tm}_{0,j} \ulcorner \Gamma \urcorner A \rightarrow \text{Tm}_{0,j} \Gamma (\text{Stage } A) \end{aligned}$$

Note that we can stage open types and terms as long as their contexts are purely object-level. By *closed staging* we mean staging only for closed types and terms.

Definition 4.2. A staging algorithm Stage is **correct** if the following properties hold:

- *Soundness:* $\ulcorner \text{Stage } A \urcorner = A$ and $\ulcorner \text{Stage } t \urcorner = t$.
- *Completeness:* staging respects definitional equality.
- *Stability:* $\text{Stage } \ulcorner A \urcorner = A$ and $\text{Stage } \ulcorner t \urcorner = t$.

We make some remarks on correctness. First, we get completeness for free in our algebraic setup, since *all* functions definable in the metatheory must respect definitional equality.

Second, note that soundness and stability together is the statement that embedding is invertible on types and terms. This is a statement of *conservativity* of 2LTT over the object theory. In [?] a significantly weaker conservativity theorem is shown, which expresses that there exists a function from $\text{Tm}_{0,j} \ulcorner \Gamma \urcorner \ulcorner A \urcorner$ to $\text{Tm}_{0,j} \Gamma A$.

Lastly, we note that the correctness terminology is borrowed from previous works on normalization-by-evaluation, where the soundness-completeness-stability trio has been used several times [?]. In [?], stability is instead called “idempotence”.

4.1 The Presheaf Model

We observe that the presheaf model described in [1, Section 2.5.3] immediately yields a closed staging algorithm, by evaluation of 2LTT types and terms in the model. In this model, contexts are presheaves over the syntactic category of the object theory. We call the model $\hat{\mathcal{O}}$ and denote its components by putting hats on 2LTT components. We summarize the key components in the following.

Notation 6. In this section, we switch to naming elements of Con_0 as a , b and c , and elements of Sub_0 as f , g , and h , to avoid name clashing with contexts and substitutions in the presheaf model.

4.1.1 The syntactic category and the meta-level fragment.

Definition 4.3. $\widehat{\mathbf{Con}} : \mathbf{Set}_{\omega+1}$ is defined as the set of presheaves over \mathbb{O} . A $\Gamma : \widehat{\mathbf{Con}}$ has an action on objects $|\Gamma| : \mathbf{Con}_{\mathbb{O}} \rightarrow \mathbf{Set}_{\omega}$ and an action on morphisms $-[-] : |\Gamma| b \rightarrow \mathbf{Sub}_{\mathbb{O}} a b \rightarrow |\Gamma| a$, such that $\gamma[\text{id}] = \gamma$ and $\gamma[f \circ g] = \gamma[f][g]$.

Notation 7. We reuse the substitution notation $-[-]$ for the action on morphisms. Also, we use lowercase γ and δ to denote elements of $|\Gamma| a$ and $|\Delta| a$ respectively.

Definition 4.4. $\widehat{\mathbf{Sub}} \Gamma \Delta : \mathbf{Set}_{\omega}$ is the set of natural transformations from Γ to Δ . A $\sigma : \widehat{\mathbf{Sub}} \Gamma \Delta$, has an action on objects $|\sigma| : \{a : \mathbf{Con}_{\mathbb{O}}\} \rightarrow |\Gamma| a \rightarrow |\Delta| a$ such that $|\sigma|(\gamma[f]) = (|\sigma| \gamma)[f]$.

Definition 4.5. $\widehat{\mathbf{T}}\mathbf{y}_{1,j} \Gamma : \mathbf{Set}_{\omega}$ is the set of *displayed presheaves* over Γ ; see e.g. [?]. This is equivalent to the set of presheaves over the category of elements of Γ , but it is usually more convenient in calculations. An $A : \widehat{\mathbf{T}}\mathbf{y} \Gamma$ has an action on objects $|A| : \{a : \mathbf{Con}_{\mathbb{O}}\} \rightarrow |\Gamma| a \rightarrow \mathbf{Set}_j$ and an action on morphisms $-[-] : |A| \gamma \rightarrow (f : \mathbf{Sub}_{\mathbb{O}} a b) \rightarrow |A|(\gamma[f])$, such that $\alpha[\text{id}] = \alpha$ and $\alpha[f \circ g] = \alpha[f][g]$.

Notation 8. We write α and β respectively for elements of $|A| \gamma$ and $|B| \gamma$.

Definition 4.6. $\widehat{\mathbf{Tm}}_{1,j} \Gamma A : \mathbf{Set}_{\omega}$ is the set of sections of the displayed presheaf A . This can be viewed as a dependently typed analogue of a natural transformation. A $t : \widehat{\mathbf{Tm}}_{1,j} \Gamma A$ has an action on objects $|t| : \{a\} \rightarrow (\gamma : |\Gamma| a) \rightarrow |A| \gamma$, such that $|t|(\gamma[f]) = (|t| \gamma)[f]$.

Using the above definitions, we can model the syntactic category of 2LTT, and also the meta-level family structure and all meta-level type formers. For an exposition in previous literature, see [?] and [?].

4.1.2 The object-level fragment. We move on to modeling the object-level syntactic fragment of 2LTT. We make some preliminary definitions. First, note the types in the object theory yield a presheaf, and terms yield a displayed presheaf over them; this immediately follows from the specification of a family structure in a cwf. Hence, we do a bit of a name overloading, and have $\mathbf{Ty}_{\mathbb{O},j} : \widehat{\mathbf{Con}}$ and $\mathbf{Tm}_{\mathbb{O},j} : \widehat{\mathbf{T}}\mathbf{y} \mathbf{Ty}_{\mathbb{O},j}$.

Definition 4.7. $\widehat{\mathbf{T}}\mathbf{y}_{\mathbb{O},j} \Gamma : \mathbf{Set}_{\omega}$ is defined as $\widehat{\mathbf{Sub}} \Gamma \mathbf{Ty}_{\mathbb{O},j}$, and $\widehat{\mathbf{Tm}}_{\mathbb{O},j} \Gamma A : \mathbf{Set}_{\omega}$ is defined as $\widehat{\mathbf{Tm}} \Gamma (\mathbf{Tm}_{\mathbb{O},j}[A])$.

For illustration, if $A : \widehat{\mathbf{T}}\mathbf{y}_{\mathbb{O},j} \Gamma$, then $A : \widehat{\mathbf{Sub}} \Gamma \mathbf{Ty}_{\mathbb{O},j}$, so $|A| : \{a : \mathbf{Con}_{\mathbb{O}}\} \rightarrow |\Gamma| a \rightarrow \mathbf{Ty}_{\mathbb{O},j} a$. In other words, the action of A on objects maps a semantic context to a syntactic object-level type. Likewise, for $t : \widehat{\mathbf{Tm}}_{\mathbb{O},j} \Gamma A$, we have $|t| : (\gamma : |\Gamma| a) \rightarrow \mathbf{Tm}_{\mathbb{O},j} a (|A| \gamma)$, so we get a syntactic object-level term as output.

Using the above definitions and following [?], we can model all type formers in $\widehat{\mathbf{T}}\mathbf{y}_{\mathbb{O},j}$. Intuitively, that is because $\widehat{\mathbf{T}}\mathbf{y}_{\mathbb{O},j}$ and $\widehat{\mathbf{Tm}}_{\mathbb{O},j}$ return types and terms, so we can reuse the types and terms in the object theory.

4.1.3 Lifting.

Definition 4.8. $\widehat{\mathbf{\Uparrow}} A : \widehat{\mathbf{T}}\mathbf{y}_{1,j} \Gamma$ is defined as $\mathbf{Tm}_{\mathbb{O},j}[A]$. With this, we get that $\widehat{\mathbf{Tm}}_{\mathbb{O},j} \Gamma A$ is equal to $\widehat{\mathbf{Tm}}_{1,j} \Gamma (\widehat{\mathbf{\Uparrow}} A)$. Hence, we can define both quoting and splicing as identity functions in the model.

4.2 Staging by Evaluation

Definition 4.9. The **evaluation morphism**, denoted \mathbb{E} is a model morphism from the syntax of 2LTT to $\hat{\mathcal{O}}$, arising from the initiality of the syntax. In other words, \mathbb{E} is defined by recursion on the syntax and strictly preserves all structure.

Note that $\hat{\cdot}$ is defined as the terminal presheaf, which is constantly \top . This leads us to the following definition.

Definition 4.10. **Closed staging** is defined as follows.

$$\begin{aligned} \text{Stage} : \text{Ty}_{0,j} \bullet &\rightarrow \text{Ty}_{0,j} \bullet & \text{Stage} : \text{Tm}_{0,j} \bullet A &\rightarrow \text{Tm}_{0,j} \bullet (\text{Stage } A) \\ \text{Stage } A &:= |\mathbb{E} A| \{ \bullet \} \text{tt} & \text{Stage } t &:= |\mathbb{E} t| \{ \bullet \} \text{tt} \end{aligned}$$

This is well-typed, since $|\mathbb{E} A| \{ \bullet \} : |\mathbb{E} \bullet| \bullet \rightarrow \text{Ty}_{0,j} \bullet$, so $|\mathbb{E} A| \{ \bullet \} : \top \rightarrow \text{Ty}_{0,j} \bullet$, and $|\mathbb{E} t| \{ \bullet \} : \top \rightarrow \text{Tm}_{0,j} \bullet (|\mathbb{E} A| \text{tt})$.

What about general (open) staging though? Given $A : \text{Ty}_{0,j} \ulcorner \Gamma \urcorner$, we get $|\mathbb{E} A| \{ \Gamma \} : |\ulcorner \Gamma \urcorner| \Gamma \rightarrow \text{Ty}_{0,j}$. We need a semantic context with type $|\ulcorner \Gamma \urcorner| \Gamma$, in order to extract an object-level type. It turns out that such “generic” semantic contexts fall out from a proof of stability for \mathbb{E} .

4.3 Stability and Open Staging

Stability means that staging an object-level construction does nothing. However, it turns out that the induction motive for contexts is necessarily the existence of generic semantic contexts. We define a $-^P$ family of functions by induction on object syntax. The induction motives are as follows.

$$\begin{aligned} (\Gamma : \text{Con}_{\mathcal{O}})^P &: |\mathbb{E} \ulcorner \Gamma \urcorner| \Gamma \\ (\sigma : \text{Sub}_{\mathcal{O}} \Gamma \Delta)^P &: \Delta^P [\sigma] = |\mathbb{E} \ulcorner \sigma \urcorner| \Gamma^P \\ (A : \text{Ty}_{0,j} \Gamma)^P &: A = |\mathbb{E} \ulcorner A \urcorner| \Gamma^P \\ (t : \text{Tm}_{0,j} \Gamma A)^P &: t = |\mathbb{E} \ulcorner t \urcorner| \Gamma^P \end{aligned}$$

The $-^P$ interpretation of all structure is straightforward. In particular, we do not have to show preservation of any definitional equality, since types, terms and substitutions are all interpreted as proof-irrelevant equations.

TODO example

Definition 4.11. We define **open staging** as follows.

$$\begin{aligned} \text{Stage} : \text{Ty}_{0,j} \ulcorner \Gamma \urcorner &\rightarrow \text{Ty}_{0,j} \Gamma & \text{Stage} : \text{Tm}_{0,j} \ulcorner \Gamma \urcorner A &\rightarrow \text{Tm}_{0,j} \Gamma (\text{Stage } A) \\ \text{Stage } A &:= |\mathbb{E} A| \Gamma^P & \text{Stage } t &:= |\mathbb{E} t| \Gamma^P \end{aligned}$$

THEOREM 4.12. *Open staging is stable.*

PROOF. For $A : \text{Ty}_{0,j}$, $\text{Stage} \ulcorner A \urcorner$ is by definition $|\mathbb{E} \ulcorner A \urcorner| \Gamma^P$, hence by A^P it is equal to A . Likewise, $\text{Stage} \ulcorner t \urcorner$ is equal to t by t^P . \square

4.4 Algorithm Extraction and Efficiency

5 YONEDA, REPRESENTABILITY AND GENERATIVITY

6 SOUNDNESS OF STAGING

REFERENCES

- [1] Danil Annenkov, Paolo Capriotti, Nicolai Kraus, and Christian Sattler. 2019. Two-Level Type Theory and Applications. *ArXiv e-prints* (may 2019). <http://arxiv.org/abs/1705.03307>

- [2] Paolo Capriotti. 2017. Models of type theory with strict equality. *arXiv preprint arXiv:1702.04912* (2017).
- [3] Vladimir Voevodsky. 2013. A simple type system with two identity types. (2013). Unpublished note.